Dynamic Mode Decomposition and Koopman theory

Sourya Dey
Galois, Inc.

Last updated on Sep 2, 2022

The material in this document extensively references [1, 2].

1 Dynamic Mode Decomposition

Say there is a vector $\boldsymbol{x}_k$, which is a vector of observed values sampled at time index k. Similarly, the vector of observed values sampled at time instant $k+1$ is $\boldsymbol{x}_{k+1}$. $\boldsymbol{x}$ has a finite number of dimensions equal to the number of observed values (say n). This means that it can be expressed as a linear combination of n basis vectors.

The transition function to take $\boldsymbol{x}_k$ to $\boldsymbol{x}_{k+1}$ is F, i.e.:

$$\boldsymbol{x}_{k+1} = F(\boldsymbol{x}_k) \tag{1}$$

$F(\cdot)$ is a non-linear function, i.e. $F(a\boldsymbol{x}_1 + b\boldsymbol{x}_2) \neq aF(\boldsymbol{x}_1) + bF(\boldsymbol{x}_2)$.[1] **Our goal is to find a matrix $\boldsymbol{A}$ such that the transition dynamics become linear**, i.e.:

$$\underset{n\times 1}{\boldsymbol{x}_{k+1}} \approx \underset{n\times n}{\boldsymbol{A}} \ \underset{n\times 1}{\boldsymbol{x}_k} \tag{2}$$

This is extremely powerful. If this indeed holds, then we can use the eigendecomposition of $\boldsymbol{A}$ to get any state of $\boldsymbol{x}$. Say the eigenvectors are $\boldsymbol{W}$ and the eigenvalues $\boldsymbol{\Lambda}$, i.e. $\underset{n\times n}{\boldsymbol{A}} \underset{n\times n}{\boldsymbol{W}} = \underset{n\times n}{\boldsymbol{W}} \underset{n\times n}{\boldsymbol{\Lambda}}$. These **eigenvectors are called the DMD modes** and they form a basis in n-dimensional space, such that any vector x can be expressed as $\boldsymbol{x} = \sum_{i=1}^{n} b_i \boldsymbol{w}_i$, where the b values are the coefficients of the linear combination. Putting them together in a vector $\boldsymbol{b}$, we get $\boldsymbol{x} = \boldsymbol{W}\boldsymbol{b}$.

Now say we have a system with known initial condition $\boldsymbol{x}_0 = \boldsymbol{W}\boldsymbol{b}$. We can solve for $\boldsymbol{b}$ as:

$$\underset{n\times 1}{\boldsymbol{b}} = \underset{n\times n}{\boldsymbol{W}^{-1}} \underset{n\times 1}{\boldsymbol{x}_0} \tag{3}$$

Then, we get $\boldsymbol{x}_1 = \boldsymbol{A}\boldsymbol{W}\boldsymbol{b} = \boldsymbol{W}\boldsymbol{\Lambda}\boldsymbol{b}$. Extending this to any time, we get the incredibly

---

[1]The function $F(\cdot)$ is written using a capital letter because typically the space is described as $\frac{d\boldsymbol{x}}{dt} = f(\boldsymbol{x})$. Integrating this from $t = k$ to $t = k+1$ gives $\boldsymbol{x}_{k+1} = F(\boldsymbol{x}_k)$.

powerful:

$$\underset{n\times 1}{\boldsymbol{x}_k} = \underset{n\times n}{\boldsymbol{W}}\ \underset{n\times n}{\boldsymbol{\Lambda}^k}\ \underset{n\times 1}{\boldsymbol{b}} = \sum_{i=1}^{n} \lambda_i{}^k b_i \underset{n\times 1}{\boldsymbol{w}_i} \tag{4}$$

for any index k way out into the past or future. Note, however, that this means that **for a stable system, the eigenvalue with largest real part (i.e. dominant eigenvalue) should have real part equal to** 1. If it exceeds 1, the system will explode in the future and implode in the past. If it's less than 1, the system will explode in the past and implode in the future.

Note that ideally *all eigenvalues* (not just the dominant one) should lie on the unit circle, because they all affect the system (albeit to a lesser extent than the dominant eigenvalue). If the non-dominant eigenvalues are close to the origin (i.e. real part much less than 1), then the system still has a chance of exploding in the past. Such numerical instability should be kept in mind when predicting states of the system in the past.

2   Implementing DMD

More generally, we have measured $\boldsymbol{x}$ at $m+1$ different time instances to get $\boldsymbol{x}_0$ to $\boldsymbol{x}_m$. These are collected as $n\times m$ matrices $\boldsymbol{X} = [\boldsymbol{x}_0, \cdots, \boldsymbol{x}_{m-1}]$ and $\boldsymbol{X}' = [\boldsymbol{x}_1, \cdots, \boldsymbol{x}_m]$. Then we get the matrix form of (2):

$$\underset{n\times m}{\boldsymbol{X}'} \approx \underset{n\times n}{\boldsymbol{A}}\ \underset{n\times m}{\boldsymbol{X}} \tag{5}$$

There can be 2 different cases:

Case 1: $m \geq n$

This is the standard regression case where we have more observations than features. We cannot solve for $\boldsymbol{A}$ exactly since we have mn equations and only n^2 variables. So, we try to minimize the least squares error $\|\boldsymbol{X}' - \boldsymbol{A}\boldsymbol{X}\|^2$, which gives the solution:

$$\underset{n\times n}{\boldsymbol{A}} = \underset{n\times m}{\boldsymbol{X}'}\ \underset{m\times n}{\boldsymbol{X}^{\dagger}} \tag{6}$$

Case 2: $m < n$

This is an over-determined system where we are trying to solve for n^2 variables in only mn equations. There can be infinite solutions. Whatever solution for A we pick, it will be an $n \times n$ matrix with rank only m.

SVD: Consider the SVD of $\boldsymbol{X}$:

$$\underset{n\times m}{\boldsymbol{X}} = \underset{n\times n}{\boldsymbol{U}}\ \underset{n\times m}{\boldsymbol{\Sigma}}\ \underset{m\times m}{\boldsymbol{V}^H} \tag{7}$$

where ^{H} is the conjugate transpose or Hermitian.

Truncated SVD: Since $m < n$, the rank of $\boldsymbol{X}$ is at most m. Let's assume the rank of $\boldsymbol{X}$ is actually r, which is less than m. (If it turns out that the rank of $\boldsymbol{X}$ is actually m, replace r with m in the following). Now, consider the truncated SVD. Since the rank is r, we need only keep the left r columns of both $\boldsymbol{U}$ and $\boldsymbol{V}$, and the top-left $r \times r$ portion of $\boldsymbol{\Sigma}$. Then we get:

$$\underset{n \times m}{\boldsymbol{X}} = \underset{n \times r}{\boldsymbol{U}} \underset{r \times r}{\boldsymbol{\Sigma}} \underset{r \times m}{\boldsymbol{V}^H} \tag{8}$$

The left singular vectors $\boldsymbol{U}$ are the proper orthogonal decomposition (POD) modes.

Reduced dimensionality POD space: The original n-dimensional vectors $\boldsymbol{x}$ can be projected into the reduced r-dimensional space of POD modes using $\tilde{\boldsymbol{x}} = \boldsymbol{U}^H \boldsymbol{x}$. Doing this to the whole matrix at once gives:

$$\underset{r \times m}{\tilde{\boldsymbol{X}}} = \underset{r \times n}{\boldsymbol{U}^H} \underset{n \times m}{\boldsymbol{X}} \tag{9}$$

and to convert back, we use $\boldsymbol{x} = \boldsymbol{U}\tilde{\boldsymbol{x}}$ and $\boldsymbol{X} = \boldsymbol{U}\tilde{\boldsymbol{X}}$.

Since r has replaced n and $m \geq r$, we are back in Case 1. We will now proceed with this case.

2.1 The DMD algorithm

Finding $\boldsymbol{A}$: Using the SVD of $\boldsymbol{X}$, we can get $\boldsymbol{A} = \boldsymbol{X}'\boldsymbol{X}^{\dagger} = \boldsymbol{X}'\boldsymbol{V}\boldsymbol{\Sigma}^{-1}\boldsymbol{U}^H$.

We can also get $\tilde{\boldsymbol{A}} = \tilde{\boldsymbol{X}}'\tilde{\boldsymbol{X}}^{\dagger}$. This is where it gets tricky. $\tilde{\boldsymbol{X}}'$ is the projection of $\boldsymbol{X}'$ on the same POD modes obtained from the SVD of $\boldsymbol{X}$, i.e. $\tilde{\boldsymbol{X}}' = \boldsymbol{U}^H \boldsymbol{X}'$. So we get:

$$\underset{r \times r}{\tilde{\boldsymbol{A}}} = \underset{r \times m}{\tilde{\boldsymbol{X}}'} \underset{m \times r}{\tilde{\boldsymbol{X}}^{\dagger}} \tag{10a}$$

$$= \underset{r \times n}{\boldsymbol{U}^H} \underset{n \times m}{\boldsymbol{X}'} \underset{m \times n}{\boldsymbol{X}^{\dagger}} \underset{n \times r}{\boldsymbol{U}} \tag{10b}$$

$$= \underset{r \times n}{\boldsymbol{U}^H} \underset{n \times n}{\boldsymbol{A}} \underset{n \times r}{\boldsymbol{U}} \tag{10c}$$

$$= \underset{r \times n}{\boldsymbol{U}^H} \underset{n \times m}{\boldsymbol{X}'} \underset{m \times r}{\boldsymbol{V}} \underset{r \times r}{\boldsymbol{\Sigma}^{-1}} \tag{10d}$$

This means that we never have to compute $\boldsymbol{A}$, we can just work with the reduced dimensionality $\tilde{\boldsymbol{A}}$. The reduced dimension r can be chosen to be anything as long as it's upper-bounded by $\min(n,m)$.

Eigendecomposition We can now compute the eigendecomposition:

$$\underset{r \times r}{\tilde{\boldsymbol{A}}} = \underset{r \times r}{\tilde{\boldsymbol{W}}} \underset{r \times r}{\tilde{\boldsymbol{\Lambda}}} \underset{r \times r}{\tilde{\boldsymbol{W}}^{-1}} \tag{11}$$

The first r eigenvalues of $\boldsymbol{A}$ are the same as those of $\tilde{\boldsymbol{A}}$, i.e. $\boldsymbol{\Lambda} = \tilde{\boldsymbol{\Lambda}}$. Since the rank of $\boldsymbol{X}$ is r (and let's assume the rank of $\boldsymbol{X}'$ is also r), the rank of $\boldsymbol{A}$ is also r (since multiplying 2 rank r matrices makes the result rank r). So, it is a good argument that the other $n - r$ eigenvalues of $\boldsymbol{A}$ are all 0.

This means that we only care about the 1st r eigenvectors of $\boldsymbol{A}$. The eigenvectors $\boldsymbol{W}$ of $\boldsymbol{A}$ can be found in 2 ways:

The *exact* eigenvectors are:

$$\underset{n \times r}{\boldsymbol{W}} = \underset{n \times m}{\boldsymbol{X}'} \underset{m \times r}{\boldsymbol{V}} \underset{r \times r}{\boldsymbol{\Sigma}^{-1}} \underset{r \times r}{\tilde{\boldsymbol{W}}} \tag{12}$$

For a backward proof, do $\boldsymbol{W}\boldsymbol{\Lambda}\boldsymbol{W}^{\dagger}$ and we will get back $\boldsymbol{A}$.

The *projected* eigenvectors are simply:

$$\underset{n \times r}{\boldsymbol{W}} = \underset{n \times r}{\boldsymbol{U}} \underset{r \times r}{\tilde{\boldsymbol{W}}} \tag{13}$$

which is what one would expect given the rules of the POD space are such that $\boldsymbol{X} = \boldsymbol{U}\tilde{\boldsymbol{X}}$. Note that this $\boldsymbol{W}$ is equal to $\boldsymbol{X}\boldsymbol{V}\boldsymbol{\Sigma}^{-1}\tilde{\boldsymbol{W}}$, which is equal to the exact formulation except that $\boldsymbol{X}'$ is replaced with $\boldsymbol{X}$. Thus, the projected will be equal to the exact if $\boldsymbol{X}$ and $\boldsymbol{X}'$ have the same column spaces.

Note that computing the projected eigenvectors is more numerically stable. However, it is recommended to use the exact formulation unless a particular eigenvalue and its corresponding exact eigenvector are both 0, in which case the projected eigenvector should be used.

Calculating other $\boldsymbol{x}$ values: The rest of the process is mostly the same. We find:

$$\underset{r \times 1}{\boldsymbol{b}} = \underset{r \times n}{\boldsymbol{W}^{\dagger}} \underset{n \times 1}{\boldsymbol{x}_0} \tag{14}$$

Then we can find past and future values as:

$$\underset{n \times 1}{\boldsymbol{x}_k} = \underset{n \times r}{\boldsymbol{W}} \underset{r \times r}{\boldsymbol{\Lambda}^k} \underset{r \times 1}{\boldsymbol{b}} = \sum_{i=1}^{r} \lambda_i{}^k b_i \underset{n \times 1}{\boldsymbol{w}_i} \tag{15}$$

2.2 Converting between discrete and continuous time

Thus far we have dealt with discrete time, i.e. the vector $\boldsymbol{x}$ was sampled at time instants indexed by k. What if we want to find a continuous time representation $\boldsymbol{x}(t)$? This is useful for, say, finding the value of $\boldsymbol{x}$ at time 3.5. In this case, the dynamics of the system are described as:

$$\underset{n \times 1}{\frac{d\boldsymbol{x}}{dt}} = \underset{n \times n}{\boldsymbol{\mathcal{A}}} \underset{n \times 1}{\boldsymbol{x}} \tag{16}$$

This has solution:

$$\boldsymbol{x}(t) = e^{\boldsymbol{\mathcal{A}}t}\boldsymbol{x}_0 \tag{17}$$

The eigendecomposition is $\boldsymbol{\mathcal{A}} = \boldsymbol{W}\boldsymbol{\Omega}\boldsymbol{W}^{-1}$. Recall that any function of $\boldsymbol{\mathcal{A}}$ gets applied to the eigenvalues, leaving the eigenvectors intact. Also, as usual, $\boldsymbol{x}_0 = \boldsymbol{W}\boldsymbol{b}$.

4

Then we can solve for any time instant as:

$$\underset{n\times 1}{\boldsymbol{x}(t)} = \boldsymbol{W}e^{\boldsymbol{\Omega}t}\boldsymbol{W}^{-1}\boldsymbol{x}_0 = \underset{n\times n}{\boldsymbol{W}}\underset{n\times n}{e^{\boldsymbol{\Omega}t}}\underset{n\times 1}{\boldsymbol{b}} \tag{18}$$

or, for the truncated case:

$$\underset{n\times 1}{\tilde{\boldsymbol{x}}(t)} = \underset{n\times r}{\boldsymbol{W}}\underset{r\times r}{e^{\boldsymbol{\Omega}t}}\underset{r\times 1}{\boldsymbol{b}} \tag{19}$$

This entire thing is equivalent to the discrete time case if we write $\boldsymbol{A} = e^{\boldsymbol{A}\Delta t}$, where Δt is the sampling interval. **This means that the eigenvectors of both matrices are the same $\boldsymbol{W}$, and the eigenvalues are related as:**

$$\lambda = e^{\omega\Delta t} \tag{20a}$$

$$\Rightarrow \omega = \frac{\ln\lambda}{\Delta t} \tag{20b}$$

As an example, say we want to find $\boldsymbol{x}$ at time 3.5. For the continuous time case, we would get $\boldsymbol{x}(3.5) = \boldsymbol{W}e^{3.5\boldsymbol{\Omega}}\boldsymbol{b}$. For the usual discrete sampling case with $\Delta t = 1$, we cannot find $\boldsymbol{x}_{3.5}$ since the index 3.5 isn't a valid sampling index. However, we can solve for $\boldsymbol{x}$ at time 3.5 if we take $\Delta t = 0.5$. **This means that we need to retake the measurements so that $\boldsymbol{X}'$, which stands for 1 index ahead of $\boldsymbol{X}$, will now be a time of 0.5 ahead.** Then, $\boldsymbol{\Lambda} = e^{0.5\boldsymbol{\Omega}}$, so $e^{3.5\boldsymbol{\Omega}} = \boldsymbol{\Lambda}^7$. This makes sense since 3.5 is the 7th index in this new sampling scheme. Then we can solve $\boldsymbol{x}(3.5) = \boldsymbol{x}_7 = \boldsymbol{W}\boldsymbol{\Lambda}^7\boldsymbol{b}$.

Tying it together: Discrete time is obviously the only way we can operate on a computer. So let's say we have measurements at times $[0, 0.5, 1, 1.5, 2, 3, 4, 5]$, and we want to predict at times 2.5 and 8. We can do it in one out of 2 ways:

- Construct the discrete problem with $\Delta t = 0.5$. This is not possible since the values at 2.5, 3.5 and 4.5 are not given. It becomes possible if we only consider times $[0, 0.5, 1, 1.5, 2]$, which, in this new scheme, become indexes $[0, 1, 2, 3, 4]$. Construct $\boldsymbol{X} = [\boldsymbol{x}_0 \cdots \boldsymbol{x}_3]$ (which is actually $[\boldsymbol{x}(0) \cdots \boldsymbol{x}(1.5)]$), $\boldsymbol{X}' = [\boldsymbol{x}_1 \cdots \boldsymbol{x}_4]$ (which is actually $[\boldsymbol{x}(0.5) \cdots \boldsymbol{x}(2)]$), and proceed normally. Once fitted, calculate $\boldsymbol{x}(2.5)$ as $\boldsymbol{x}_5$ and $\boldsymbol{x}(8)$ as $\boldsymbol{x}_{16}$ using (15).

- Construct the discrete problem with $\Delta t = 1$. Only consider times $[0, 1, 2, 3, 4, 5]$. Construct $\boldsymbol{X} = [\boldsymbol{x}_0 \cdots \boldsymbol{x}_4]$, $\boldsymbol{X}' = [\boldsymbol{x}_1 \cdots \boldsymbol{x}_5]$, and proceed normally. Once fitted, convert the discrete eigenvalues $\boldsymbol{\Lambda}$ to continuous eigenvalues $\boldsymbol{\Omega}$ using (20), and calculate $\boldsymbol{x}(2.5)$ and $\boldsymbol{x}(5)$ using (19).

2.3   Different starting index

Thus far we have assumed that the starting index is $\boldsymbol{x}_0$. What if it's not? Suppose we have measurements starting at $\boldsymbol{x}_i$, $i \neq 0$. There are 2 methods to deal with this:

Assume index 0 exists: Assume there is an un-measured $\boldsymbol{x}_0$. Then, $\boldsymbol{x}_i = \boldsymbol{A}^i \boldsymbol{x}_0$. Putting the earlier expression $\boldsymbol{x}_0 = \boldsymbol{W}\boldsymbol{b}$ into this yields:

$$
\begin{aligned}
\boldsymbol{x}_i &= \left(\boldsymbol{W}\boldsymbol{\Lambda}\boldsymbol{W}^{-1}\right)^i \boldsymbol{W}\boldsymbol{b} \\
&= \boldsymbol{W}\boldsymbol{\Lambda}^i \boldsymbol{b} \\
\Rightarrow \boldsymbol{b} &= \boldsymbol{\Lambda}^{-i}\boldsymbol{W}^{-1}\boldsymbol{x}_i
\end{aligned}
\tag{21}
$$

In terms of the specific notation used so far, (14) now becomes:

$$
\underset{r\times 1}{\boldsymbol{b}} = \underset{r\times r}{\boldsymbol{\Lambda}^{-i}}\,\underset{r\times n}{\boldsymbol{W}^{\dagger}}\,\underset{n\times 1}{\boldsymbol{x}_i}
\tag{22}
$$

Thus, we can get $\boldsymbol{b}$ in terms of whichever index $\boldsymbol{x}$ starts from by using the appropriate opposite power on the eigenvalue matrix. We don't need to know $\boldsymbol{x}_0$.

Then, any other value $\boldsymbol{x}_j$ can be obtained as:

$$
\begin{aligned}
\boldsymbol{x}_j &= \boldsymbol{W}\boldsymbol{\Lambda}^j \boldsymbol{b} \\
&= \boldsymbol{W}\boldsymbol{\Lambda}^j \boldsymbol{\Lambda}^{-i}\boldsymbol{W}^{\dagger}\boldsymbol{x}_i \\
&= \boldsymbol{A}^{j-i}\boldsymbol{x}_i
\end{aligned}
\tag{23}
$$

which is as one would expect when computing index j from index i. Also note that j can be less than i, and everything will still hold. For example, $\boldsymbol{x}_0 = \boldsymbol{A}^{-i}\boldsymbol{x}_i$.

Shift the indexes so that i becomes 0: Thus just means subtracting i from any index. So, the given $\boldsymbol{x}_i$ in the original index space becomes $\boldsymbol{x}_{i-i} = \boldsymbol{x}_0$ in the shifted index space, and $\boldsymbol{x}_j$ becomes $\boldsymbol{x}_{j-i}$. That way, we retain the relation $\boldsymbol{b} = \boldsymbol{W}^{-1}\boldsymbol{x}_0$ in the shifted index space, i.e. we avoid powers on the eigenvalue matrix when computing $\boldsymbol{b}$ and can keep the math the same.

Now, suppose we want to compute $\boldsymbol{x}_j$ in the original index space. Then, in the shifted index space, we need to compute:

$$
\boldsymbol{x}_{j-i} = \boldsymbol{W}\boldsymbol{\Lambda}^{j-i}\boldsymbol{b}
\tag{24}
$$

as one would expect. Also note that $\boldsymbol{x}_0$ in the original index space will be computed in the shifted index space as:

$$
\begin{aligned}
\boldsymbol{x}_{0-i} &= \boldsymbol{W}\boldsymbol{\Lambda}^{0-i}\boldsymbol{b} \\
&= \boldsymbol{W}\boldsymbol{\Lambda}^{0-i}\boldsymbol{W}^{-1}\boldsymbol{x}_0 \\
&= \boldsymbol{A}^{-i}\boldsymbol{x}_0 \\
&= \boldsymbol{A}^{-i}\boldsymbol{x}_{i-i}
\end{aligned}
$$

When expressed in the original index space, this becomes $\boldsymbol{x}_0 = \boldsymbol{A}^{-i}\boldsymbol{x}_i$, which is identical to the previous method. Thus, **the two methods give identical results, as they should**.

2.4   Multiple trajectories

A trajectory is defined as a run of a dynamical system from some initial state to some final state, i.e. $[\boldsymbol{x}_0, \cdots, \boldsymbol{x}_m]$. So far, we have been dealing with single trajectories. But

it can so happen that a dynamical system is described by multiple trajectories instead of a single trajectory. For example, consider the system $\boldsymbol{x}_{k+1} = \boldsymbol{x}_k + \psi$, where ψ denotes noise. We want a single $\boldsymbol{A}$ for the whole system since its dynamics don't change. But we also need a way to represent data from the multiple trajetories. This can be done by *concatenating the trajectories along the index dimension.*

Suppose we perform J runs of a system, which have $[m_1 + 1, m_2 + 1, \cdots, m_J + 1]$ indexed states, respectively. Using superscript for trajectory, the states can be described as $\left[\boldsymbol{x}_0^1, \cdots, \boldsymbol{x}_{m_1}^1\right], \left[\boldsymbol{x}_0^2, \cdots, \boldsymbol{x}_{m_2}^2\right], \cdots, \left[\boldsymbol{x}_0^J, \cdots, \boldsymbol{x}_{m_J}^J\right]$. Concatenating along the index dimension, we get:

$$\boldsymbol{X} = \begin{bmatrix} \boldsymbol{x}_0^1, \cdots, \boldsymbol{x}_{m_1}^1, & \boldsymbol{x}_0^2, \cdots, \boldsymbol{x}_{m_2}^2, & \cdots, & \boldsymbol{x}_0^J, \cdots, \boldsymbol{x}_{m_J-1}^J \end{bmatrix} \tag{25a}$$

$$\boldsymbol{X}' = \begin{bmatrix} \boldsymbol{x}_1^1, \cdots, \boldsymbol{x}_{m_1}^1, & \boldsymbol{x}_0^2, \cdots, \boldsymbol{x}_{m_2}^2, & \cdots, & \boldsymbol{x}_0^J, \cdots, \boldsymbol{x}_{m_J}^J \end{bmatrix} \tag{25b}$$

Both these matrices have dimensions $n \times (M - 1)$, where $M = \sum_{j=1}^J m_j$. $\boldsymbol{X}$ leaves out the last sample of the last trajectory, while $\boldsymbol{X}'$ leaves out the first sample of the first trajectory. This means that for $\boldsymbol{X}$, $M - 1$ is parsed as $m_1 + m_2 + \cdots + (m_J - 1)$, while for $\boldsymbol{X}'$, the same $M - 1$ is parsed as $(m_1 - 1) + m_2 + \cdots + m_J$.

The rest of the algorithm proceeds as given previously, except that m is replaced by $M - 1$. This means that the reduced dimension (i.e. rank) r should be upper-bounded by $\min(n, M - 1)$.

3 Koopman theory

The problem with DMD is that the original dynamics may not be linearizable. However, what if we transform the problem from the existing n-dimensional space to a different space?

The Koopman operator replaces all vectors $\boldsymbol{x}$ with functions $g(\boldsymbol{x})$. The output of $g(\boldsymbol{x})$ can be a scalar or vector. Essentially, now the transformed value(s) at time k is $g(\boldsymbol{x}_k)$ and at time $k + 1$ is $g(\boldsymbol{x}_{k+1})$. Since $g(\boldsymbol{x})$ is a function, in general, it has an infinite number of dimensions, i.e. the number of basis functions whose linear combination equals $g(\boldsymbol{x})$ is infinite.

The transition operator to take g($\boldsymbol{x}_k$) to g($\boldsymbol{x}_{k+1}$) is K, which is the Koopman operator. So:

$$g(\boldsymbol{x}_{k+1}) = Kg(\boldsymbol{x}_k) \tag{26}$$

We assume that K is a linear operator, i.e. it is a (infinite-dimensional) matrix. **Essentially, the Koopman operator transforms the problem of a non-linear transition function acting in a finite-dimensional vector space to the problem of a linear operator acting on an infinite-dimensional function space.**

The basis functions for this transformed space can be the eigenfunctions of K, i.e. the functions $\{\phi(\cdot)\}$ which satisfy $K\phi(\boldsymbol{x}) = \lambda\phi(\boldsymbol{x})$. So, any function g on which K operates can be written as a linear, infinite combination of the Koopman eigenfunctions $\{\phi(\cdot)\}$.

3.1 Koopman operator as function composition

Note that $Kg(\boldsymbol{x}_k) = g(\boldsymbol{x}_{k+1})$. But as per (1), $\boldsymbol{x}_{k+1} = F(\boldsymbol{x}_k)$. This implies that $Kg(\boldsymbol{x}_k) = g(F(\boldsymbol{x}_k))$. Dropping the index:

$$Kg(\boldsymbol{x}) = g(F(\boldsymbol{x})) \tag{27}$$

So, the Koopman operator can be regarded as function composition.

3.2 Combining with Dynamic Mode Decomposition

Let us assume the Koopman space is finite and p-dimensional. Then, we can write:

$$\underset{p \times 1}{\boldsymbol{y}} = g\left(\underset{n \times 1}{\boldsymbol{x}}\right) \tag{28}$$

Note that the term *extended dynamic mode decomposition (EDMD)* refers to this same technique, except it uses orthonormal polynomial basis functions. Koopman theory in more general and can work with any $g(\cdot)$. One can use a neural network to represent $g(\cdot)$.

Once we have the $\boldsymbol{y}$ values, **one can perform DMD exactly as described thus far by replacing $\boldsymbol{x}$ with $\boldsymbol{y}$, $\boldsymbol{X}$ with $\boldsymbol{Y}$, n with p, and solving for $\boldsymbol{y}_k$.** Everything else remains the same, so if $p > m$, one can work with $r < p$ and so on. Finally, one needs to convert back using:

$$\underset{n \times 1}{\boldsymbol{x}} = h\left(\underset{p \times 1}{\boldsymbol{y}}\right) \tag{29}$$

where $h(\cdot)$ should approximate $g^{-1}(\cdot)$.

References

[1] J. N. Kutz, S. L. Brunton, B. W. Brunton, and J. L. Proctor. *Dynamic Mode Decomposition: Data Driven Modeling of Complex Systems*. Society for Industrial and Applied Mathematics, 2016.

[2] Jonathan H. Tu, Clarence W. Rowley, Dirk M. Luchtenburg, Steven L. Brunton, and J. Nathan Kutz. On dynamic mode decomposition: Theory and applications. *Journal of Computational Dynamics*, 1(2):391–421, 2014.

Appendix: Variable reference

m – Number of data measurements excluding the 0th measurement. Thus, total number of data measurements is $m + 1$.

n – Number of input states.

p – Number of encoded states.

r – Rank.